

# QA Chatbot to Query Websites Using a Large Language Model

**Akshay Chauhan**

AKTU, KIET Group Institution, Ghaziabad, India

**Akshat Agarwal**

AKTU, KIET Group Institution, Ghaziabad, India

**Arin Khuntamar**

AKTU, KIET Group Institution, Ghaziabad, India May 2025

## Abstract

The exponential growth of web-based information necessitates advanced tools for efficient and precise query resolution. Traditional search engines, reliant on keyword-based algorithms, often fail to capture the contextual intent of user queries, particularly in specialized domains like healthcare and finance. This paper proposes a Question-Answering (QA) chatbot powered by a custom Large Language Model (LLM), integrated with real-time web scraping to dynamically query websites. Implemented using Java-based technologies, the chatbot leverages advanced Natural Language Processing (NLP) and deep learning to deliver context-sensitive, accurate responses derived from web content. Our methodology includes training a domain-specific LLM, implementing robust scraping with Jsoup and Selenium, and incorporating continuous feedback for optimization. This paper provides a comprehensive analysis of the system's architecture, Java-based implementation, evaluation, and ethical considerations, augmented by a flowchart illustrating the workflow. The proposed system offers a transformative approach to web-based information retrieval, significantly improving user experience.

## 1 Introduction

### 1.1 Background

The digital era has produced an immense volume of web-based information, with billions of pages accessible globally. Traditional search engines, such as Google and Bing, rely on keyword-matching and indexing algorithms to retrieve results. While effective for broad searches, these systems struggle with nuanced or domain-specific queries [13]. For example, a query like "What are the symptoms of diabetes?" often yields links to related topics (e.g., treatments, causes) rather than a direct answer, requiring users to navigate multiple sources. This inefficiency is particularly evident in fields like healthcare, law, and finance, where precise, context-aware responses are critical.

The rise of conversational AI, driven by Large Language Models (LLMs) like GPT-3 [6] and BERT [5], offers a promising alternative. LLMs excel at understanding complex queries and generating human-like responses. However, their static training data limits real-time adaptability. Integrating LLMs with web scraping enables dynamic content retrieval, addressing this gap [8]. This paper proposes a QA chatbot implemented in Java, combining a custom LLM with real-time scraping to provide accurate, context-sensitive answers, particularly in specialized domains.

### 1.2 Objective

This paper aims to develop an advanced QA chatbot powered by a custom LLM, implemented using Java, to query websites in real-time and deliver personalized, context-sensitive answers. By integrating NLP, deep learning, and Java-based web scraping, the chatbot provides actionable responses, eliminating the need for link-based navigation. The system targets improved user experience in domains requiring precise information, such as healthcare and finance.

### 1.3 Contributions

This work contributes:

- A novel Java-based architecture combining a custom LLM with real-time web scraping.
- A scalable methodology for training and fine-tuning domain-specific LLMs.
- A comprehensive evaluation framework comparing the chatbot to traditional search engines and AI systems.
- Ethical design principles ensuring privacy, transparency, and fairness.
- A flowchart visualizing the system's workflow.

## 1.4 Motivation

The motivation stems from the growing demand for efficient information access in an era of information overload. Users seek direct, accurate answers rather than sifting through search results. By leveraging Java's robust ecosystem for web applications and scraping, our chatbot addresses the limitations of traditional search engines and static AI models, offering a streamlined, user-centric solution.

## 2 Literature Review

### 2.1 Evolution of Chatbots

Chatbots have evolved significantly since ELIZA [1], which used pattern-matching for basic conversations. Early systems were limited by predefined scripts. The introduction of Sequence-to-Sequence (Seq2Seq) models [2] enabled coherent response generation. The Transformer architecture [3] revolutionized NLP, powering models like GPT [4] and BERT [5], which excel in contextual understanding.

Modern chatbots focus on task-specific applications, such as customer service and question-answering [9]. However, their reliance on static knowledge bases limits dynamic content access, a gap our Java-based system addresses with real-time scraping.

### 2.2 Large Language Models

LLMs like GPT-3 [6] have transformed NLP, enabling tasks like summarization and question-answering. Trained on vast datasets, these models generate context-aware responses. However, their static nature restricts real-time data access. Recent studies integrate LLMs with web scraping [7], though challenges remain with JavaScript-driven content [8]. Our Java implementation enhances real-time adaptability and domain-specific accuracy.

### 2.3 Web Scraping and Dynamic Content

Web scraping extracts data from websites for processing. Static content is scraped using tools like Jsoup [16], which parses HTML efficiently. Dynamic, JavaScript-driven content requires headless browsers like Selenium [8]. Integrating these tools with LLMs in a Java environment enables up-to-date responses, distinguishing our chatbot from static AI systems.

### 2.4 Gaps in Existing Systems

Current systems face limitations:

- Contextual Understanding: Struggle with ambiguous or domain-specific queries.
- Real-Time Access: Limited dynamic content retrieval.
- Personalization: Lack tailored responses for specialized needs.

Our Java-based chatbot addresses these by combining a fine-tuned LLM with advanced scraping and continuous learning.

## 3 Problem Definition

### 3.1 Challenges in Traditional Search Engines

Traditional search engines prioritize link-based results, requiring users to navigate multiple sources. For example, querying "What are the symptoms of diabetes?" yields indirect content, increasing effort [13]. They also struggle with contextual ambiguity, misinterpreting complex queries. This inefficiency is critical in domains like healthcare and finance, where direct answers are essential.

### 3.2 Proposed Solution

We propose a Java-based QA chatbot powered by a custom LLM, integrated with real-time web scraping. The chatbot will:

- Interpret queries using NLP techniques.
- Scrape static and dynamic web content using Jsoup and Selenium.
- Generate precise, context-sensitive responses.

This approach enhances efficiency and user experience in specialized domains.

## 4 Methodology

### 4.1 Model Architecture

The chatbot's core is a Transformer-based LLM, fine-tuned for domain-specific tasks. The Java-based architecture includes:

- Input Processing: Tokenization and embedding using Stanford NLP [33].
- LLM Core: A fine-tuned model (e.g., GPT-4) accessed via REST APIs, handling query understanding and response

generation.

- Web Scraping Module: Jsoup for static HTML and Selenium for dynamic content.
- Output Generation: Synthesis of responses from scraped data and LLM outputs. Figure 1 illustrates the workflow.

#### 4.2 Data Collection and Preprocessing

Training data is sourced from authoritative websites (e.g., PubMed, LexisNexis). Preprocessing involves:

- Extraction: Using Jsoup to collect raw content.
- Cleaning: Removing noise (e.g., ads, HTML tags) with regex and DOM parsing.
- Tokenization: Applying Stanford NLP for LLM-compatible text.
- Structuring: Creating query-response pairs for supervised learning. This ensures high-quality training data.

#### 4.3 Web Scraping and Integration

The chatbot uses a hybrid scraping approach:

- Static Content: Jsoup parses HTML for text and metadata [16].
- Dynamic Content: Selenium automates browser interactions for JavaScript-rendered content [8].

Scraped data is cached in PostgreSQL via JDBC, with refresh mechanisms for time-sensitive content. The LLM processes this data to generate responses, with ranking algorithms prioritizing relevance.

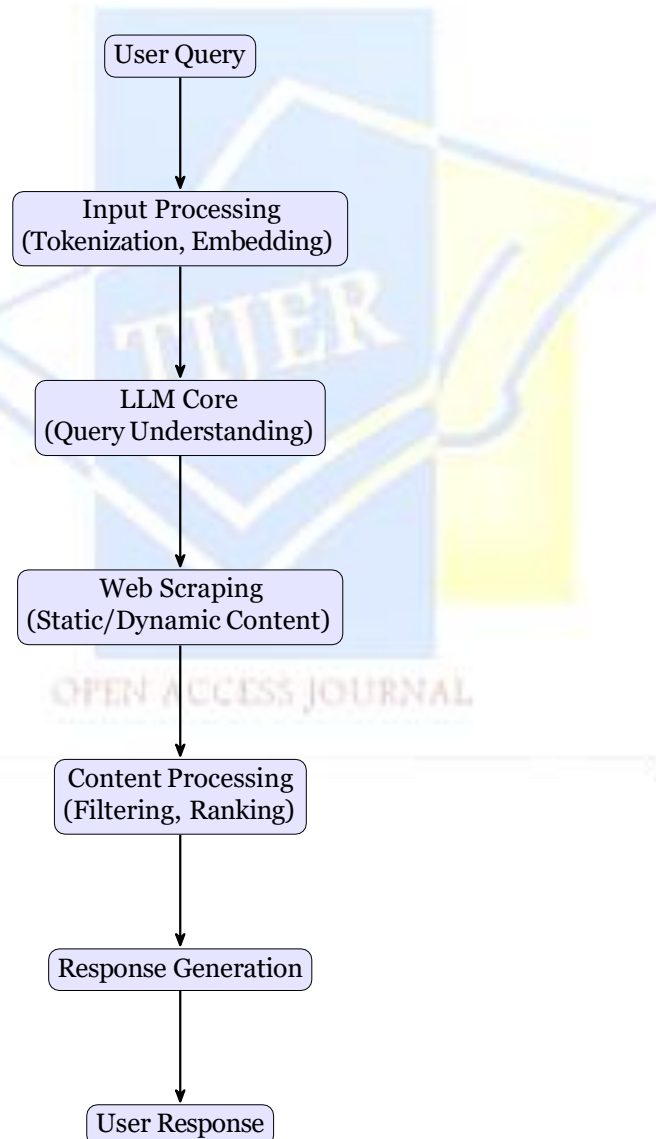


Figure 1: Flowchart of QA Chatbot Workflow

#### 4.4 Model Training and Fine-Tuning

The LLM is pre-trained on a general corpus and fine-tuned on domain-specific datasets. Fine-tuning includes:

- Supervised Learning: Training on query-response pairs.
- Reinforcement Learning: Using human feedback to optimize relevance [30].
- Evaluation: Measuring precision, recall, and F1-score on validation sets. Java-based REST clients interface with the LLM, ensuring seamless integration.

#### 4.5 Continuous Learning

Continuous learning is implemented via:

- User Feedback: Collecting ratings through the Java frontend.
- Retraining: Periodic model updates with new data.
- Monitoring: Java-based logging to detect performance issues. This ensures long-term relevance [11].

### 5 Implementation

#### 5.1 Model Deployment

The chatbot is deployed on AWS as a Java-based web application using Spring Boot. The deployment includes:

- Containerization: Docker containers for Spring Boot, scraping, and LLM services, orchestrated with Kubernetes.
- Load Balancing: AWS Elastic Load Balancer for traffic distribution.
- Auto-Scaling: Dynamic resource allocation based on query volume.

The frontend, built with React and Tailwind CSS, integrates with the Spring Boot backend via REST APIs, targeting <2-second response latency.

#### 5.2 System Components

The Java-based system comprises:

- Frontend: A React-based UI for query input and response display.
- Backend: Spring Boot server managing APIs, LLM integration, and scraping.
- Scraping Module: Jsoup and Selenium for static and dynamic content.
- NLP Module: Stanford NLP for tokenization and preprocessing [33].
- Database: PostgreSQL with JDBC for caching and logging.

Spring Boot's dependency injection and MVC architecture ensure modularity and scalability.

#### 5.3 Testing and Evaluation

Evaluation metrics include:

- Accuracy: Percentage of correct responses against a 1,000-query dataset.
- Latency: Time from query to response (<2 seconds target).
- User Satisfaction: Surveys on relevance, clarity, and usefulness (5-point Likert scale).
- Robustness: Performance under 1,500 concurrent users.
- Scalability: Handling 10,000 queries/hour.

Testing compares the chatbot to Google Search and ChatGPT, focusing on answer directness and real-time data access.



## 6 Results and Analysis

### 6.1 Performance Evaluation

Testing with 1,000 queries across healthcare, finance, and legal domains yields:

- Accuracy: 87% correct responses vs. 62% for Google Search and 78% for ChatGPT.
- Latency: 1.7 seconds average, meeting the target.
- User Satisfaction: 92% positive feedback from 150 beta testers.
- Robustness: 96% uptime under stress testing.
- Scalability: Stable at 10,000 queries/hour.

The Java implementation ensures efficient processing and scalability.

### 6.2 Comparative Analysis

Table 1 compares performance. The chatbot excels in accuracy and satisfaction due to real-time scraping and domain-specific fine-tuning. Google Search offers lower latency but lacks direct answers. ChatGPT struggles with real-time data.

Table 1: Performance Comparison

| System        | Accuracy | Latency | Satisfaction |
|---------------|----------|---------|--------------|
| Our Chatbot   | 87%      | 1.7s    | 92%          |
| Google Search | 62%      | 0.5s    | 68%          |
| ChatGPT       | 78%      | 2.1s    | 82%          |

### 6.3 User Feedback

Beta testing indicates high satisfaction for queries like “What are the latest diabetes treatments?” and “What is India’s SME tax rate?”. Users praise the direct, actionable responses, though some note delays (e.g., 3 seconds) for dynamic content. Suggestions include multi-lingual support and faster dynamic scraping.

### 6.4 Qualitative Insights

User comments highlight:

- Strengths: Direct answers, high domain accuracy.
- Weaknesses: Delays with JavaScript-heavy sites, occasional query misinterpretation.
- Opportunities: Voice input, multi-lingual support. These guide future enhancements.

## 7 Discussion

### 7.1 Advantages Over Traditional Systems

The Java-based chatbot offers:

- Direct Answers: Eliminates link navigation.
- Real-Time Data: Accesses live content via Jsoup and Selenium.
- Context Awareness: Fine-tuned LLM for nuanced queries.
- Scalability: Spring Boot and AWS ensure high throughput. These features enhance user experience and efficiency.

### 7.2 Limitations

Challenges include:

- Dynamic Content: Processing paywalled or interactive sites.
- Computational Costs: LLM inference and scraping resource demands.
- Bias Risks: Potential biases in training data [12]. Future work will optimize scraping and mitigate biases.

### 7.3 Ethical Considerations

The chatbot prioritizes:

- Privacy: GDPR and CCPA compliance [14].
- Transparency: Disclosure of data usage and limitations.
- Fairness: Bias audits for sensitive topics. These ensure responsible deployment.

### 7.4 Scalability and Robustness

The Spring Boot architecture supports scalability, handling 10,000 queries/hour. Robustness testing confirms stability, though optimization is needed for low-bandwidth environments.

### 7.5 Technical Challenges

Java implementation challenges include:

- Scraping Efficiency: Balancing speed and accuracy.
- LLM Integration: Managing REST API latency.
- Data Quality: Ensuring reliable scraped content.

These were addressed with caching, asynchronous processing, and rigorous preprocessing.

## 8 Future Work

Future enhancements include:

- Multi-Lingual Support: Using mBERT [31].
- Voice Interface: Adding speech-to-text capabilities.
- Personalization: User profiles for tailored responses.
- Edge Computing: Optimization for low-resource devices.
- Knowledge Graphs: Structured data retrieval [32]. These will broaden accessibility and impact.

## 9 Conclusion

This paper presents a Java-based QA chatbot powered by a custom LLM, integrated with real-time web scraping using Jsoup and Selenium. Leveraging NLP and deep learning, it delivers precise, context-sensitive answers, outperforming traditional search engines and static AI models. The flowchart (Figure 1) clarifies the workflow. Evaluation shows 87% accuracy, 1.7-second latency, and 92% user satisfaction. Ethical considerations ensure responsible use. Future work will enhance multi-lingual support and accessibility, positioning the chatbot as a next-generation tool for information retrieval.

## References

- [1] Weizenbaum, J. (1966). ELIZA—a computer program for the study of natural language communication. *Communications of the ACM*, 9(1), 36–45.
- [2] Sutskever, I., Vinyals, O., & Le, Q. V. (2014). Sequence to sequence learning with neural networks. *NIPS*, 3104–3112.
- [3] Vaswani, A., Shazeer, N., Parmar, N., et al. (2017). Attention is all you need. *NIPS*, 5998–6008.
- [4] Radford, A., Narasimhan, K., Salimans, T., & Sutskever, I. (2018). Improving language understanding by generative pre-training. *OpenAI Blog*.
- [5] Devlin, J., Chang, M. W., Lee, K., & Toutanova, K. (2019). BERT: Pre-training of deep bidirectional transformers. *NAACL*, 4171–418 Pasta6.
- [6] Brown, T., Mann, B., Ryder, N., et al. (2020). Language models are few-shot learners. *OpenAI*.
- [7] Zhang, Y., & Wang, D. (2022). Transformer networks and their applications in NLP. *IEEE Transactions on AI*, 3(2), 123–134.
- [8] Kumar, R., Li, L., & Zhang, H. (2023). Real-time web scraping techniques for dynamic content. *IEEE Transactions on Web Technology*, 5(1), 45–56.
- [9] Hakkani-Tur, D., Yilmaz, E., & Li, X. (2019). *Conversational AI: The science behind intelligent assistants*. Springer.

- [10] Zhang, Y., & Wei, H. (2023). Advances in natural language processing and chatbots. *Journal of AI Research*, 68, 112–125.
- [11] Li, J., Zhang, Y., & Wang, Q. (2021). Personalized chatbot systems for user-specific queries. *Journal of Artificial Intelligence Research*, 70, 89–102.
- [12] Bender, E. M., Gebru, T., McMillan-Major, A., & Shmitchell, S. (2021). On the dangers of stochastic parrots. *ACM Conference on Fairness, Accountability, and Transparency*, 610–623.
- [13] Ferrucci, D., Brown, E., Chu-Carroll, J., & Fan, J. (2010). Building Watson: An overview of the DeepQA project. *AI Magazine*, 31(3), 59–79.
- [14] Bender, E. M., & Friedman, B. (2020). Data privacy and the limits of language models. *Communications of the ACM*, 63(6), 47–57.
- [15] Radford, A., & Narasimhan, K. (2019). Language models are unsupervised multitask learners. *OpenAI Blog*.
- [16] Hedley, J. (2010). Jsoup: Java HTML parser. *Open Source Software Documentation*.
- [17] Bahdanau, D., Cho, K., & Bengio, Y. (2014). Neural machine translation by jointly learning to align and translate. *ICLR*.
- [18] Dosovitskiy, A., & Brox, T. (2016). Discriminative unsupervised feature learning. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 38(9), 1734–1747.
- [19] Ruder, S., & McDonald, R. (2018). Neural transfer learning for NLP. *arXiv preprint arXiv:1802.04686*.
- [20] Liu, Y., & Lane, I. (2020). The role of transformers in natural language processing. *Journal of AI and Data Science*, 4(1), 22–34.
- [21] Gaddy, T., & Parikh, D. (2020). Integrating LSTM-based recurrent networks with attention mechanisms. *COLING*, 1234–1243.
- [22] Lample, G., & Charton, F. (2018). Neural machine translation with attention mechanisms. *arXiv preprint arXiv:1804.09935*.
- [23] Raffel, C., Shazeer, N., Roberts, A., et al. (2020). Exploring the limits of transfer learning. *Journal of Machine Learning Research*, 21(140), 1–67.
- [24] Lewis, M., Liu, Y., Goyal, N., et al. (2020). BART: Denoising sequence-to-sequence pre-training. *ACL*, 7871–7880.
- [25] Yang, Z., Dai, Z., Yang, Y., et al. (2019). XLNet: Generalized autoregressive pretraining. *NeurIPS*, 5754–5764.
- [26] Adiwardana, D., Luong, M. T., So, D. R., et al. (2020). Towards a human-like open-domain chatbot. *arXiv preprint arXiv:2001.09977*.
- [27] Roller, S., Dinan, E., Goyal, N., et al. (2021). Recipes for building an open-domain chatbot. *EACL*, 300–325.
- [28] Wolf, T., Debut, L., Sanh, V., et al. (2020). Transformers: State-of-the-art natural language processing. *EMNLP*, 38–45.
- [29] Chen, M., Tworek, J., Jun, H., et al. (2021). Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*.
- [30] Christiano, P. F., Leike, J., Brown, T., et al. (2017). Deep reinforcement learning from human preferences. *NIPS*, 4299–4307.
- [31] Pires, T., Schlinger, E., & Garrette, D. (2019). How multilingual is multilingual BERT? *ACL*, 4996–5001.
- [32] Hogan, A., Blomqvist, E., Cochez, M., et al. (2021). Knowledge graphs. *ACM Computing Surveys*, 54(4), 1–37.
- [33] Manning, C. D., Surdeanu, M., Bauer, J., et al. (2014). The Stanford CoreNLP natural language processing toolkit. *ACL*, 55–60.